

# Operating systems

Processes

Lecture 2

Florenc Demrozi

[florenc.demrozi@univr.it](mailto:florenc.demrozi@univr.it)

University of Verona  
Department of Computer Science

2021/2022



# Table of Contents

## 1 Attributes of a process

- Identifier
- Environment
- Working directory
- File descriptor table

## 2 Operations with processes

- Termination
- Creation
- Monitoring

## 3 Program execution (exec library functions)



# Attributes of a process



# Recall to the previous lecture

## Process

A *process* is an instance of an executing program.

From the Kernel's point of view, a process consists of:

- user-space memory containing program code,
- the variables used by that code, and
- a set of kernel data structures that maintain information about the process's state (e.g. page tables, table of open files, signals to be delivered, process resource usage and limits, ...)



# Attributes of a process

## Identifier



# Process identifier

The `getpid` system call returns the process ID of the calling process.

```
#include <unistd.h>
#include <sys/types.h>

pid_t getpid(void);
```

The `pid_t` data type used for the return value of `getpid` is an integer type for the purpose of storing process IDs.

With the exception of a few system processes such as `init` (process ID 1), there is no fixed relationship between a program and the process ID of the process that is created to run that program.

```
// to see the init process
user@localhost[~]$ ps auxf
```

N.B. The `getpid` system call is always successful!



## Real and effective process user-ID (1/3)

The `getuid` and `getgid` system calls return, respectively, the real user ID and real group ID of the calling process. The `geteuid` and `getegid` system calls perform the corresponding tasks for the effective IDs.

```
#include <unistd.h>
#include <sys/types.h>

uid_t getuid(void); // Real user ID
uid_t geteuid(void); // Effective user ID

gid_t getgid(void); // Real group ID
gid_t getegid(void); // Effective group ID
```

N.B. They are always successful!



## Real and effective process user-ID (2/3)

- **real** user ID and group ID identify the user and group to which the process belongs,
- **effective** user ID and group ID are used to determine the permissions granted to a process when it tries to perform operations.

Here is the content of file `program.c`:

```
#include <unistd.h>
#include <sys/types.h>
int main (int argc, char *argv[]) {
    printf("PID: %d, user-ID: real %d, effective %d\n",
        getpid(), getuid(), geteuid());
    return 0;
}
```





## Real and effective process user-ID (3/3)

```
user@localhost[~]$ gcc -o program program.c
user@localhost[~]$ ls -l program
-r-xr-xr-x 1 Professor Professor 8712 Jan 16 16:27 program

user@localhost[~]$ ./program
PID: 1234, user-ID: real 1000, effective 1000

user@localhost[~]$ sudo ./program
PID: 1423, user-ID: real 0, effective 0

user@localhost[~]$ sudo chmod u+s program
user@localhost[~]$ ls -l program
-r-sr-xr-x 1 root Professor 8712 Jan 16 16:27 program

user@localhost[~]$ ./program
PID: 4321, user-ID: real 1000, effective 0
```

Keep in mind: if the **S**ticky bit is unset, then the user's permissions are granted to the executable to perform operations. If it is set, then the owner's permissions are granted to executable.



# Attributes of a process


## Environment



# Process environment (1/5)

Each process has an associated array of strings called the environment list, or simply the environment. Each of these strings is a definition of the form `name = value`. When a new process is created, it **inherits** a copy of its parent's environment.

The structure of the environment list is as follows:



## Process environment (2/5)

Within a C program, the environment list can be accessed by either:

- 1 using the global variable `char **environ`. Originally it was used specifically in POSIX systems, now this technique is widely used and supported by many systems.
- 2 or you can also receive the current environment as third argument of the `main` function. This technique is recognized as standard C, but it is not supported by all the operating systems.



## Process environment (3/5)

Displaying the process environment, **first** technique:

```
#include <stdio.h>
// Global variable pointing to the environment of the process.
extern char **environ;

int main(int argc, char *argv[]) {
    for (char **it = environ; (*it) != NULL; ++it) {
        printf("--> %s\n", *it);
    }
    return 0;
}
```

```
user@localhost[~]$ ./program
--> $HOME=/home/Professor
--> $PWD=/tmp
--> $USER=Professor
```



## Process environment (4/5)

Displaying the process environment, **second** technique:

```
#include <stdio.h>
int main(int argc, char *argv[], char* env[]) {
    for (char **it = env; (*it) != NULL; ++it) {
        printf("--> %s\n", *it);
    }
    return 0;
}
```

```
user@localhost[~]$ ./program
--> $HOME=/home/Professor
--> $PWD=/tmp
--> $USER=Professor
```



## Process environment (5/5)

```
#include <stdlib.h>
// Returns pointer to (value) string, or NULL if no such variable exists
char *getenv(const char *name);
// Returns 0 on success, or -1 on error
int setenv(const char *name, const char *value, int overwrite);
// Returns 0 on success, or -1 on error
int unsetenv(const char *name);
```

- given a variable name, `getenv` returns a pointer to its string value, or `NULL` if no environment variable exists with the specified name.
- `setenv` adds `name=value` to the environment, unless a variable identified by `name` already exists and `overwrite` has the value 0. If `overwrite` is nonzero, the environment is always changed.
- `unsetenv` removes the variable identified by `name` from the environment.



# Attributes of a process

Working directory





# Process working directory (1/3)

A process can retrieve its current working directory using `getcwd`.

```
#include <unistd.h>

// Returns cwdbuf on success, or NULL on error.
char *getcwd(char *cwdbuf, size_t size);
```

On success, `getcwd` returns a pointer to `cwdbuf` as its function result. If the pathname for the current working directory exceeds `size` bytes, then `getcwd` returns `NULL`.

The caller must allocate the `cwdbuf` buffer to be at least `size` bytes in length. (Normally, we would size `cwdbuf` using the `PATH_MAX` constant.)



## Process working directory (2/3)

The `chdir` system call changes the calling process's current working directory to the relative or absolute pathname specified in `pathname`.

```
#include <unistd.h>

// Returns 0 on success, or -1 on error
int chdir(const char *pathname);
```

The `fchdir` system call does the same as `chdir`, except that the directory is specified via a file descriptor previously obtained by opening the directory with `open`.

```
#define _BSD_SOURCE
#include <unistd.h>

// Returns 0 on success, or -1 on error.
int fchdir(int fd);
```



# Process working directory (3/3)

```
char buf[PATH_MAX];  
// Open the current working directory  
int fd = open(".", O_RDONLY);  
getcwd(buf, PATH_MAX);  
printf("1) Current dir:\n\t%s\n", buf);  
  
// Move the process into /tmp  
chdir("/tmp");  
getcwd(buf, PATH_MAX);  
printf("2) Current dir:\n\t%s\n", buf);  
  
// Move the process back into the initial directory  
fchdir(fd);  
getcwd(buf, PATH_MAX);  
printf("3) Current dir:\n\t%s\n", buf);  
  
// Close the file descriptor  
close(fd);
```

Here is the output:

```
1) Current dir:  
    /home/Professor  
2) Current dir:  
    /tmp  
3) Current dir:  
    /home/Professor
```



# Attributes of a process

## File descriptor table



## Process file descriptor table (1/5)

Each process has an associated *file descriptor table*. Each entry represents an input/output resource (e.g. file, pipe, socket) used by the process. The directory `/proc/<PID>/fd1` contains a symbolic link for each entry of the file descriptor table of a process.

A created process has always three file descriptors (`stdin`, `stdout`, `stderr`)

```
user@localhost[~]$ sleep 30 &
[1] 1344

user@localhost[~]$ ls -l /proc/1344/fd
total 0
lrwx----- 1 Professor Professor 0 Jan 18 12:35 0 -> /dev/pts/0
lrwx----- 1 Professor Professor 0 Jan 18 12:35 1 -> /dev/pts/0
lrwx----- 1 Professor Professor 0 Jan 18 12:35 2 -> /dev/pts/0
```

---

<sup>1</sup>Process information pseudo-file system, it doesn't contain 'real' files but runtime system information.



# Process file descriptor table (2/5)

## Displaying the file descriptor entries of a process

```
char buf[PATH_MAX];  
// Replace %i with PID, and store the resulting string in buf.  
snprintf(buf, PATH_MAX, "/proc/%i/fd/", getpid());  
  
DIR *dir = opendir(buf);  
struct dirent *dp;  
while ((dp = readdir(dir)) != NULL) {  
    if ((strcmp(dp->d_name, ".") != 0) && (strcmp(dp->d_name, "..") != 0)) {  
        printf("\tEntry: %s\n", dp->d_name);  
    }  
}  
closedir(dir);
```

```
user@localhost[~]$ ./program  
Entry: 0 // link to stdin  
Entry: 1 // link to stdout  
Entry: 2 // link to stderr  
Entry: 3 // link to /proc/<PID>/fd directory
```



## Process file descriptor table (3/5)

### Important

At a new entry of the file descriptor table is always assigned the lowest available index.

Redirecting the standard output stream of a process to a file named myfile

```
// We close STDOUT which has FD 1. The remaining file descriptors have
// index 0 (stdin) and 2 (stderr).
close(STDOUT_FILENO);
// We open a new file, to which will be assigned FD 1 automatically
// because it is the lowest available index in the table.
int fd = open("myfile", O_TRUNC | O_CREAT | O_WRONLY, S_IRUSR | S_IWUSR);
// Printf uses the FD 1, thus, it will print on the file.
printf("ciao\n");
```

No string will be displayed on terminal, as stdout stream is closed.  
However, all string printed by printf will be reported in myfile.



## Process file descriptor table (4/5)

The `dup` system call takes an open file descriptor, and returns a new descriptor that refers to the same open file description. The new descriptor is guaranteed to be the lowest unused file descriptor.

```
#include <unistd.h>

// Returns (new) file descriptor on success, or -1 on error.
int dup(int oldfd);
```





# Process file descriptor table (5/5)

Redirecting stdout and stderr of a process to a file named myfile

```
// FDT: [0, 1, 2] -> [0, 2]
close(STDOUT_FILENO);
// FDT: [0, 2]      -> [0]
close(STDERR_FILENO);
// FDT: [0]         -> [0, 1]
int fd = open("myfile", O_TRUNC | O_CREAT | O_WRONLY, S_IRUSR | S_IWUSR);
// FDT: [0, 1]      -> [0, 1, 2]
dup(1);
// FDT: [0: STDIN, 1: myfile, 2: myfile]
printf("Have a good ");
fflush(stdout);
fprintf(stderr, "day!\n");
```

```
user@localhost[~]$ cat myfile
Have a good day!
```



# Operations with processes



# Operations with processes

## Termination



# Process termination (1/5)

The process calling `_exit()` is always successfully terminated.

```
#include <unistd.h>

void _exit(int status);
```

The first byte of the `status` argument defines the *termination status* of the process. By convention, the zero value indicates that the process terminated **successfully**, a nonzero status value indicates that the process terminated **unsuccessfully**.



## Process termination (2/5)

Programs generally call `exit()` rather than `_exit()`.

```
#include <stdlib.h> // N.B. provided by C library  
  
void exit(int status);
```

C library defines the macros `EXIT_SUCCESS` (0) and `EXIT_FAILURE` (1)

The following actions are performed by `exit()` method:

- Call exit handlers (see next slides).
- Flush `stdio` stream buffers.
- Call `_exit()`, using the value supplied in `status`.



## Process termination (3/5)

An *exit handler* is a function that is registered during the life of a process. It is automatically called during the process termination via *exit()*.

```
#include <stdlib.h>

// Returns 0 on success, or nonzero on error.
int atexit(void (*func)(void));
```

The `atexit()` adds the provided function pointer `func` to a list of functions that are called during the process termination.

`func` has to be defined to take no argument and return no value.

If more *exit handler* are registered, then they are called in reverse order of registration.



# Process termination (4/5)

```
#include <stdlib.h>
#include <stdio.h>
#include <unistd.h>
void func1() { printf("\tAtexit function 1 called\n"); }
void func2() { printf("\tAtexit function 2 called\n"); }
int main (int argc, char *argv[]) {
    if (atexit(func1) != 0 || atexit(func2) != 0)
        _exit(EXIT_FAILURE);
    exit(EXIT_SUCCESS);
}
```

Here is the output of the program:

```
user@localhost[~]$ ./exit_handlers
Atexit function 2 called
Atexit function 1 called
```



# Process termination (5/5)

One other way in which a process may terminate is to return from `main()`

- performing an explicit `return n` is equivalent to calling `exit(n)`;
- performing an implicit `return` or falling off the end of `main()` is equivalent to calling `exit(0)` in C99 standard. Otherwise, the behaviour of the process is undefined.





# Operations with processes

## Creation



# Process creation (1/6)

```
#include <unistd.h>

// In parent: returns process ID of child on success, or -1 on error.
// In created child: always returns 0.
pid_t fork(void);
```

The `fork()` system call creates a new process, the *child*, which is an almost exact duplicate of the calling process, the *parent*.

- After the execution of a `fork()`, two processes exist, and, in each process, execution continues from the point where the `fork()` returns.
- It is indeterminate which of the two processes is next scheduled to use the CPU.
- The *child* receives duplicates of all parent's file descriptors and the attached shared memories (see Filesystem and IPC slides)



## Process creation (2/6)

```
#include <unistd.h>
int main (int argc, char *argv[]) {
    int stack = 111;
    pid_t pid = fork();
    if (pid == -1)
        errExit("fork");
    // -->Both parent and child come here !!!<--
    if (pid == 0)
        stack = stack * 4;
    printf("\t%s stack %d\n", (pid==0) ? "(child )" : "(parent)", stack);
}
```



## Process creation (3/6)

Here is the output of the program:

```
user@localhost[~]$ ./example_fork
(parent) stack 111
(child ) stack 444

user@localhost[~]$ ./example_fork
(child ) stack 444
(parent) stack 111
```

The terminal output shows that:

- the *child* process gets its own copy of the parent's variables;
- the execution of both *parent* and *child* processes continue from the point where the `fork()` returned;



## Process creation (4/6)

Each process has a parent, namely the process that created it (see previous slides about `fork()`).

```
#include <unistd.h>

// Always successfully returns PID of caller's parent.
pid_t getppid(void);
```

The ancestor of all processes is the process `init` (PID=1). If a *child* process becomes **orphaned** because its *parent* terminates, then the *child* is “adopted” by the process `init`. The subsequent calls to `getppid()` in the *child* return 1.



# Process creation (5/6)

```
#include <unistd.h>
int main (int argc, char *argv[]) {
    pid_t pid = fork();
    if (pid == -1) {
        errExit("fork");
    }
    if (pid == 0) {
        printf("(child ) PID: %d PPID: %d\n", getpid(), getppid());
    }
    else {
        printf("(parent) PID: %d PPID: %d\n", getpid(), getppid());
    }
    return 0;
}
```



## Process creation (6/6)

The execution of the previous example has **three** different scenarios:

- 1 The *child* is executed after the *parent*, and the *parent* is not terminated

```
(parent) PID: 402 PPID: 350  
(child ) PID: 403 PPID: 402
```

- 2 The *child* is executed before the *parent* process

```
(child ) PID: 403 PPID: 402  
(parent) PID: 402 PPID: 350
```

- 3 The *child* is executed after the termination of the *parent*

```
(parent) PID: 402 PPID: 350  
(child ) PID: 403 PPID: 1
```

Question: Whose the process ID 350 may belong to?



# Operations with processes

## Monitoring





# Monitoring a child Process (1/12)

The `wait` system call waits for one of the children of the calling process to terminate. (see `waitpid` for *status* input argument).

```
#include <sys/wait.h>

// Returns PID of terminated child, or -1 on error.
pid_t wait(int *status);
```

The following actions are performed by `wait`:

- If calling process have no unwaited-for children, then `wait` returns -1 and `errno` is `ECHILD`.
- If no child has yet terminated, then `wait` blocks the calling process until a child terminates. If a child has already terminated, then `wait` returns immediately.
- If `status` is not `NULL`, information about the terminated child is stored in the integer variable which `status` points to (next slides).



# Monitoring a child Process (2/12)

```
for (int i = 1; i <= 3; ++i) {  
    // Fork and ignore fork failures.  
    if (fork() == 0) {  
        printf("Child %d sleeps %d seconds...\n", getpid(), i);  
        // Suspends the calling process for i seconds  
        sleep(i);  
        _exit(0);  
    }  
}  
  
pid_t child;  
while ((child = wait(NULL)) != -1)  
    printf("wait() returned child %d\n", child);  
if (errno != ECHILD)  
    printf("(wait) An unexpected error...\n");
```



# Monitoring a child Process (3/12)

Example output:

```
user@localhost[~]$ ./example_wait
child 75 sleeps 1 seconds
child 76 sleeps 2 seconds
child 77 sleeps 3 seconds
wait() returned child 75
wait() returned child 76
wait() returned child 77
```

**Question:** What happens to a child that terminates before its parent has had the chance to perform a `wait`?



## Monitoring a child Process (4/12)

The kernel deals with this situation by turning the terminated child into a *zombie* process. This means that most of the resources held by the child are released back to the system. The only parts of the terminated process still maintained are:

- 1 its process ID;
- 2 its termination status;
- 3 the resource usage statistics.

If the parent process terminates without calling `wait`, then the zombie child process is “adopted” by the process `init`, which will perform a `wait` system call some time later eventually.



# Monitoring a child Process (5/12)

The `waitpid` system call suspends execution of the calling process until a child specified by `pid` argument has changed state.

```
#include <sys/wait.h>

// Returns a PID, 0, or -1 on error.
pid_t waitpid(pid_t pid, int *status, int options);
```

The `status` argument is the same of `wait` (see next slides). The value of `pid` determines what child process we want to wait.

- $\text{pid} \geq 0$ , wait for the child having PID equals to `pid`.
- $\text{pid} = 0$ , wait for any child in the same caller's process group<sup>2</sup>.
- $\text{pid} < -1$ , wait for any child in the process group `|pid|`.
- $\text{pid} = -1$ , wait for any child.

---

<sup>2</sup>The processes can be organized in process group and sessions to support shell job control in Linux (topic not included in Operating System program)



# Monitoring a child Process (6/12)

```
#include <sys/wait.h>

// Returns a PID, 0, or -1 on error.
pid_t waitpid(pid_t pid, int *status, int options);
```

The options argument of the waitpid system call is an OR of zero or more of the following constants:

- WUNTRACED: return when a child is stopped by a signal or it terminates.
- WCONTINUED: return when a child has been resumed by delivery of a SIGCONT signal.
- WNOHANG: If no child specified by pid has yet changed state, then return immediately, instead of blocking (i.e., perform a “poll”). In this case, the return value of waitpid is 0.
- 0: then waitpid waits only for terminated children.



# Monitoring a child Process (7/12)

```
pid_t pid;
for (int i = 0; i < 3; ++i) {
    pid = fork();
    if (pid == 0) {
        // Code executed by the child process...
        _exit(0);
    }
}
// The parent process only waits for the last created child
waitpid(pid, NULL, 0);
```



# Monitoring a child Process (7/12)

```
pid_t pid = fork();
if (pid == 0) {
    // Code executed by the child process
} else {
    // Waiting for a terminated/stopped | resumed child process.
    waitpid(pid, NULL, WUNTRACED | WCONTINUED);
}
```






# Monitoring a child Process (9/12)

The status value set by `waitpid`, and `wait`, let us distinguish the following events for a child process:

- 1 The child process terminated by calling `_exit` (or `exit`).




- The macro `WIFEXITED` returns true if the child exited normally.
- The macro `WEXITSTATUS` returns the exit status of the child process.

```
waitpid(-1, &status, WUNTRACED | WCONTINUED);
if (WIFEXITED(status)) {
    printf("Child exited, status=%d\n", WEXITSTATUS(status));
}
```



# Monitoring a child Process (10/12)

- ② The child was terminated by the delivery of an unhandled signal.



- The macro `WIFSIGNALED` returns true if the child was killed by a signal.
- The macro `WTERMSIG` returns the number of the signal that caused the process to terminate.


```
waitpid(-1, &status, WUNTRACED | WCONTINUED);
if (WIFSIGNALED(status)) {
    printf("child killed by signal %d (%s)",
        WTERMSIG(status), strsignal(WTERMSIG(status)));
}
```

The `strsignal(int sig)` is a method of `string.h` which returns a string describing the signal `sig` (see IPC part 1).



# Monitoring a child Process (11/12)

- ③ The child was stopped by a signal.



- The macro `WIFSTOPPED` returns true if the child process was stopped by a signal.
- The macro `WSTOPSIG(status)` returns the number of the signal that stopped the process.

```
waitpid(-1, &status, WUNTRACED | WCONTINUED);
if (WIFSTOPPED(status)) {
    printf("child stopped by signal %d (%s)\n",
        WSTOPSIG(status), strsignal(WSTOPSIG(status)));
}
```



# Monitoring a child Process (12/12)

- 4 The child was resumed by a SIGCONT signal.



- The macro WIFCONTINUED returns true if the child was resumed by delivery of SIGCONT.

```
waitpid(-1, &status, WUNTRACED | WCONTINUED);  
if (WIFCONTINUED(status)) {  
    printf("child resumed by a SIGCONT signal\n");  
}
```

or

```
waitpid(-1, &status, WCONTINUED);  
printf("child resumed by a SIGCONT signal\n");
```



## Program execution (exec library functions)



# The exec Library Functions (1/2)

The exec family of functions replaces the current process image with a new process image.

```
#include <unistd.h>
// None of the following returns on success, all return -1 on error.
int execl (const char *path, const char *arg, ... ); // ... variadic functions
int execlp(const char *path, const char *arg, ... );
int execl_e(const char *path, const char *arg, ... , char *const envp[]);
int execv (const char *path, char *const argv[]);
int execvp(const char *path, char *const argv[]);
int execve(const char *path, char *const argv[], char *const envp[]);
```

**Note:** The list of arguments must be terminated by a NULL pointer, and, since these are variadic functions, this pointer must be cast (char \*) NULL.



## The exec Library Functions (2/2)

| function | path     | arguments (argv) | environment (envp) |
|----------|----------|------------------|--------------------|
| execl    | pathname | list             | caller's environ   |
| execlp   | filename | list             | caller's environ   |
| execle   | pathname | list             | array              |
| execv    | pathname | array            | caller's environ   |
| execvp   | filename | array            | caller's environ   |
| execve   | pathname | array            | array              |

- **path:** *pathname* means the absolute path to an executable. While *filename* means the name of an executable, which is sought in the list of directories specified in the PATH environment variable.
- **argv:** a NULL-terminated list/array of pointers to string defining the command line argument of the program.
- **envp:** a NULL-terminated array of pointers to string (name = value) defining the environment of the program.



# Example (1/2)

Program: example.c

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
int main(int argc, char *argv[])
{
    printf("PID of example.c = %d\n", getpid());
    char *args[] = {"Hello", "C", "Programming", NULL};
    execv("./hello", args);
    printf("Back to example.c");
    return 0;
}
```





## Example (2/2)

### Program: hello.c

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
int main(int argc, char *argv[])
{
    printf("We are in hello.c\n");
    printf("PID of hello.c = %d\n", getpid());
    return 0;
}
```

```
user@localhost[~]$ gcc -o example example.c
user@localhost[~]$ gcc -o hello hello.c
user@localhost[~]$ ./example
PID of example.c = 4733
We are in Hello.c
PID of hello.c = 4733
```



# execl(...) function

Using the program `printenv` to print the environment variable `HOME`.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    execl("/usr/bin/printenv", "printenv", "HOME", (char *)NULL);
    perror("Execl");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (pathname) | : "/usr/bin/printenv"              |
| argv (list)     | : "printenv", "HOME", (char *)NULL |
| envp (-)        | : caller's environment             |

---

```
user@localhost[~]$ ./example_exec
/home/user
```



# execvp(...) function

Using the program printenv to print the environment variable HOME.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    execvp("printenv", "printenv", "HOME", (char *)NULL);
    perror("Execp");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (filename) | : "printenv"                       |
| argv (list)     | : "printenv", "HOME", (char *)NULL |
| envp (-)        | : caller's environment             |

---

```
user@localhost[~]$ ./example_exec
/home/user
```



# execle(...) function

Using the program printenv to print the environment variable HOME.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    char *env[] = {"HOME=/home/pippo", (char *)NULL};
    execle("/usr/bin/printenv", "printenv", "HOME", (char *)NULL, env);
    perror("Execle");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (pathname) | : "/usr/bin/printenv"              |
| argv (list)     | : "printenv", "HOME", (char *)NULL |
| envp (array)    | : "HOME=/home/pippo", (char *)NULL |

---

```
user@localhost[~]$ ./example_exec
/home/pippo
```



# execv(...) function

Using the program printenv to print the environment variable HOME.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    char *arg[] = {"printenv", "HOME", (char *)NULL};
    execv("/usr/bin/printenv", arg);
    perror("Execv");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (pathname) | : "/usr/bin/printenv"              |
| argv (array)    | : "printenv", "HOME", (char *)NULL |
| envp (-)        | : caller's environment             |

---

```
user@localhost[~]$ ./example_exec
/home/user
```



# execvp(...) function

Using the program `printenv` to print the environment variable `HOME`.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    char *arg[] = {"printenv", "HOME", (char *)NULL};
    execvp("printenv", arg);
    perror("Execvp");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (filename) | : "printenv"                       |
| argv (array)    | : "printenv", "HOME", (char *)NULL |
| envp (-)        | : caller's environment             |

---

```
user@localhost[~]$ ./example_exec
/home/user
```



## execve(...) function

Using the program printenv to print the environment variable HOME.

```
#include <unistd.h>
#include <stdio.h>
int main (int argc, char *argv[]) {
    char *arg[] = {"printenv", "HOME", (char *)NULL};
    char *env[] = {"HOME=/home/pippo", (char *)NULL};
    execve("/usr/bin/printenv", arg, env);
    perror("Execve");
}
```

---

|                 |                                    |
|-----------------|------------------------------------|
| path (pathname) | : "/usr/bin/printenv"              |
| argv (array)    | : "printenv", "HOME", (char *)NULL |
| envp (array)    | : "HOME=/home/pippo", (char *)NULL |

---

```
user@localhost[~]$ ./example_exec
/home/pippo
```



# Final remarks on the exec library functions

What you should always keep in mind when you use an `exec` function:

- The program input parameter points to an executable;
- List and array are always terminated with a NULL pointer

```
(char *)NULL;
```

- By convention, the first item of `argv` is the name of the program;
- All `exec` functions return no result on success.

