

Operating systems

Interprocess communication (IPC)

Part 1 of 3: System V IPC

Semaphores

Lecture 4.1

Florenc Demrozi

florenc.demrozi@univr.it

University of Verona
Department of Computer Science

2021/2022



Table of Contents

1 Introduction to System V IPC

- Creating and Opening
- Data Structures

2 IPCs Commands

- ipcs
- ipcrm

3 Semaphores

- Creating and Opening
- Control Operations
- Other Operations



Introduction to System V IPC



Introduction to System V IPC

Unix System V (aka "System Five")

Unix System V is one of the first commercial versions of the Unix operating system. It was originally developed by AT&T and first released in 1983. Four major versions of System V were released, numbered 1, 2, 3, and 4. System V is sometimes abbreviated to SysV.

Interprocess communication (IPC)

Interprocess communication (IPC) refers to mechanisms that coordinate activities among cooperating processes. A common example of this need is managing access to a given system resource.



Introduction to System V IPC

System V IPCs refers to three different mechanisms for interprocess communication:

- **Semaphores** let processes to synchronize their actions. A semaphore is a kernel-maintained value, which is appropriately modified by system's processes before performing some critical actions
- *Message queues* can be used to pass messages among processes.
- *Shared memory* enables multiple processes to share a their region of memory.

Other IPC

- **Signals**
- Pipes
- FIFOs



Introduction to System V IPC

Creating and Opening



Creating and opening a System V IPC object

Each System V IPC mechanism has an associated *get* system call (`msgget`, `semget`, or `shmget`), which is analogous to the `open` system call.

Given an integer *key* (analogous to a filename), the *get* system call can either first create a new IPC, and then returns its unique identifier, or returns the identifier of an existing IPC.

An IPC *identifier* is analogous to a *file descriptor*. It is used in all subsequent system calls to refer to the IPC object.



Creating and opening a System V IPC object

Example showing how to create a semaphore (overview)

```
// PERM: rw-----
id = semget(key, 10 ,IPC_CREAT | S_IRUSR | S_IWUSR);
if (id == -1)
    errExit(semget);
```

As with all of the *get* calls, the *key* is the first argument. It is a value sensible for the application using the IPC object. The returned IPC *identifier* is a unique code identifying the IPC object in the system.

Mapping with the *open(...)* system call:

key -> filename

id -> file descriptor



System V IPC keys

System V IPC keys are integer values represented using the data type `key_t`. The IPC get calls translate a key into the corresponding integer IPC identifier.

So, how do we provide a unique key that guarantees we do not accidentally obtain the identifier of an existing IPC object used by some other application?



System V IPC keys - IPC_PRIVATE flag

When creating a new IPC object, the *key* may be specified as `IPC_PRIVATE`. In this way, we delegate the problem of finding a unique key to the kernel.

Example of the usage of `IPC_PRIVATE`:

```
id = semget(IPC_PRIVATE, 10, S_IRUSR | S_IWUSR);
```

This technique is especially useful in *multiprocess applications* where the parent process creates the IPC object prior to performing a `fork()`, with the result that the child inherits the identifier of the IPC object.



System V IPC keys - `ftok()`

The `ftok` (file to key) function converts a `pathname` and a `proj_id` (i.e., project identifier) to a System V IPC key.

```
#include <sys/ipc.h>

// Returns integer key on success, or -1 on error (check errno)
key_t ftok(char *pathname, int proj_id);
```

The provided `pathname` has to refer to an existing, accessible file. The last 8 bits of `proj_id` are actually used, and they have to be a nonzero value).

Typically, `pathname` refers to one of the files, or directories, created by the application.



System V IPC keys - `ftok()`

Example shows a typical usage of the function `ftok`

```
key_t key = ftok("/mydir/myfile", 'a');
if (key == -1)
    errExit("ftok failed");

int id = semget(key, 10, S_IRUSR | S_IWUSR);
if (id == -1)
    errExit("semget failed");
```

Example: Character "a"

- ASCII = 097
- Binary = 01100001



Introduction to System V IPC

Data Structures



Associated Data Structure - ipc_perm

The kernel maintains an associated data structure (`msqid_ds`, `semid_ds`, `shmid_ds`) for each instance of a System V IPC object. As well as data specific to the type of IPC object, each associated data structure **includes** the substructure `ipc_perm` holding the granted permissions.

```
struct ipc_perm {  
    key_t __key;           /* Key, as supplied to 'get' call */  
    uid_t uid;             /* Owner's user ID */  
    gid_t gid;             /* Owner's group ID */  
    uid_t cuid;            /* Creator's user ID */  
    gid_t cgid;            /* Creator's group ID */  
    unsigned short mode;   /* Permissions */  
    unsigned short __seq;  /* Sequence number */  
};
```



Associated Data Structure - ipc_perm

- The `uid` and `gid` fields specify the ownership of the IPC object.
- The `cuid` and `cgid` fields hold the user and group IDs of the process that created the object.
- The `mode` field holds the permissions mask for the IPC object, which are initialized using the lower 9 bits of the flags specified in the `get` system call used to create the object.

Some important notes about `ipc_perm`:

- 1 The `cuid` and `cgid` fields are immutable.
- 2 Only read and write permissions are meaningful for IPC objects. Execute permission is meaningless, and it is ignored.



Associated Data Structure - ipc_perm - Example

Example shows a typical usage of the `semctl` to change the owner of a semaphore.

```
struct semid_ds semq;  
// get the data structure of a semaphore from the kernel  
if (semctl(semid, 0, IPC_STAT, &semq) == -1)  
    errExit("semctl get failed");  
// change the owner of the semaphore  
semq.sem_perm.uid = newuid;  
// update the kernel copy of the data structure  
if (semctl(semid, IPC_SET, &semq) == -1)  
    errExit("semctl set failed");
```

Similarly, the `shmctl` and `msgctl` system calls are applied to update the kernel data structure of a *shared memory* and *message queue*.



IPCs Commands



IPCs Commands

ipcs



The `ipcs` command

Using `ipcs`, we can obtain information about IPC objects on the system. By default, `ipcs` displays all objects, as in the following example:

```
user@localhost[~]$ ipcs
----- Message Queues -----
key      msqid      owner      perms      used-bytes  messages
0x1235   26          student    620        12          20

----- Shared Memory Segments -----
key      shmid      owner      perms      bytes       nattch     status
0x1234   0          professor  600        8192        2

----- Semaphore Arrays -----
key      semid      owner      perms      nsems
0x1111   102       professor  330        20
```



IPCs Commands

ipcrm



The `ipcrm` command

Using `ipcrm`, we can remove IPC objects from the system.

Remove a message queue:

```
ipcrm -Q 0x1235    ( 0x1235 is the key of a queue )  
ipcrm -q 26        ( 26 is the identifier of a queue )
```

Remove a shared memory segment

```
ipcrm -M 0x1234    ( 0x1234 is the key of a shared memory seg. )  
ipcrm -m 0          ( 0 is the identifier of a shared memory seg. )
```

Remove a semaphore array

```
ipcrm -S 0x1111    ( 0x1111 is the key of a semaphore array )  
ipcrm -s 102        ( 102 is the identifier of a semaphore array )
```



Semaphores



Semaphores

Creating and Opening



Creating/Opening a Semaphore Set

The `semget` system call creates a new **semaphore set** or obtains the identifier of an existing set.

```
#include <sys/sem.h>

// Returns semaphore set identifier on success, or -1 error
int semget(key_t key, int nsems, int semflg);
```

The key arguments are: an IPC key, `nsems` specifies the number of semaphores in that set, and must be greater than 0. `semflg` is a bit mask specifying the permissions (see `open(...)` system call, `mode` argument) to be placed on a new semaphore set or checked against an existing set.

In additions, the following flags can be ORed (`|`) in `semflg`:

- `IPC_CREAT`: If no semaphore set with the specified key exists, create a new set.
- `IPC_EXCL`: in conjunction with `IPC_CREAT`, it makes `semget` fail if a semaphore set exists with the specified key.



Creating/Opening a Semaphore Set

Example showing how to create a semaphore set having 10 semaphores

```
int semid;
key_t key = //... (generate a key in some way, i.e. with ftok)

// A) delegate the problem of finding a unique key to the kernel
semid = semget(IPC_PRIVATE, 10, S_IRUSR | S_IWUSR);

// B) create a semaphore set with identifier key, if it doesn't already exist
semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);

//C) create a semaphore set with identifier key, but fail if it exists already
semid = semget(key, 10, IPC_CREAT | IPC_EXCL | S_IRUSR | S_IWUSR);
```



Semaphore Control Operations

The `semctl` system call performs a variety of control operations on a semaphore set or on an individual semaphore within a set.

```
#include <sys/sem.h>

// Returns nonnegative integer on success, or -1 error
int semctl(int semid, int semnum, int cmd, ... /* union semun arg */);
```

The `semid` argument is the identifier of the semaphore set on which the operation is to be performed.

Certain control operations (`cmd`) require a third/fourth argument. Before presenting the available control operations on a semaphore set, the union `semun` is introduced.



Semaphore Control Operations - union *semun*

The union `semun` must be **explicitly defined by the programmer** before calling the `semctl` system call.

```
#ifndef SEMUN_H
#define SEMUN_H
#include <sys/sem.h>
// definition of the union semun
union semun {
    int val;
    struct semid_ds * buf;
    unsigned short * array;
};
#endif
```



Semaphores

Control Operations



Semaphore Control Operations

Generic control operations

```
int semctl(semid, 0 /*ignored*/, cmd, arg);
```

- IPC_RMID: Immediately remove the semaphore set. Any processes blocked is awakened (error set to EIDRM). The arg argument is not required.
- IPC_STAT: Place a copy of the `semid_ds` data structure associated with this semaphore set in the buffer pointed to by `arg.buf`.
- ICP_SET: Update selected fields of the `semid_ds` data structure associated with this semaphore set using values in the buffer pointed to by `arg.buf`.



Semaphore Control Operations

Generic control operations

```
struct semid_ds {  
    struct ipc_perm sem_perm; /* Ownership and permissions */  
    time_t sem_otime;         /* Time of last semop() */  
    time_t sem_ctime;         /* Time of last change */  
    unsigned long sem_nsems; /* Number of semaphores in set */  
};
```

Only the subfields *uid*, *gid*, and *mode* of the substructure *sem_perm* can be updated via `IPC_SET`.



Semaphore Control Operations

Generic control operations (Example)

Example showing how to **change the permissions** of a semaphore set

```
ket_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// instantiate a semid_ds struct
struct semid_ds ds;
// instantiate a semun union (defined manually somewhere)
union semun arg;
arg.buf = &ds;
// get a copy of semid_ds structure belonging to the kernel
if (semctl(semid, 0 /*ignored*/, IPC_STAT, arg) == -1)
    errExit("semctl IPC_STAT failed");
// update permissions to guarantee read access to the group
arg.buf->sem_perms.mode |= S_IRGRP;
// update the semid_ds structure of the kernel
if (semctl(semid, 0 /*ignored*/, IPC_SET, arg) == -1)
    errExit("semctl IPC_SET failed");
```



Semaphore Control Operations

Generic control operations (Example)

Example showing how to **remove** semaphore set

```
if (semctl(semid, 0/*ignored*/, IPC_RMID, 0/*ignored*/) == -1)
    errExit("semctl failed");
else
    printf("semaphore set removed successfully\n");
```



Semaphore Control Operations

Retrieving and initializing semaphore values

```
int semctl(semid, semnum, cmd, arg);
```

- SETVAL: the value of the *semnum-th* semaphore in the set referred to by *semid* is initialized to the value specified in *arg.val*.
- GETVAL: as its function result, *semctl* returns the value of the *semnum-th* semaphore in the semaphore set specified by *semid*. The *arg* argument is not required.

```
int semctl(semid, 0 /*ignored*/, cmd, arg);
```

- SETALL: initialize all semaphores in the set referred to by *semid*, using the values supplied in the array pointed to by *arg.array*.
- GETALL: retrieve the values of all of the semaphores in the set referred to by *semid*, placing them in the array pointed to by *arg.array*.



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **initialize a specific semaphore** in a semaphore set

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// set the semaphore value to 0
union semun arg;
arg.val = 0;
// initialize the 5-th semaphore to 0
if (semctl(semid, 5, SETVAL, arg) == -1)
    errExit("semctl SETVAL");
```

A semaphore set must be always initialized before using it!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **get the current state** of a specific semaphore in a semaphore set.

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);

// get the current state of the 5-th semaphore
int value = semctl(semid, 5, GETVAL, 0/*ignored*/);
if (value == -1)
    errExit("semctl GETVAL");
```

Once returned, the semaphore may already have changed state!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **initialize a semaphore** set having 10 semaphores

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// set the first 5 semaphores to 1, and the remaining to 0
int values[] = {1,1,1,1,1,0,0,0,0,0};
union semun arg;
arg.array = values;
// initialize the semaphore set
if (semctl(semid, 0/*ignored*/, SETALL, arg) == -1)
    errExit("semctl SETALL");
```

A semaphore set must be always initialized before using it!



Semaphore Control Operations

Retrieving and initializing semaphore values (Example)

Example showing how to **get the current state** of a semaphore set having 10 semaphores

```
ket_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// declare an array big enough to store the semaphores' value
int values[10];
union semun arg;
arg.array = values;
// get the current state of a semaphore set
if (semctl(semid, 0/*ignored*/, GETALL, arg) == -1)
    errExit("semctl GETALL");
```

Once returned, a semaphore may already have changed state!



Semaphore Control Operations

Retrieving per-semaphore information

```
int semctl(semid, semnum, cmd, 0);
```

- GETPID: return the process ID of the last process to perform a semop on the *semnum-th* semaphore
- GETNCNT: return the number of processes currently waiting for the value of the *semnum-th* semaphore to increase
- GETZCNT: return the number of processes currently waiting for the value of the *semnum-th* semaphore to become 0;



Semaphore Control Operations

Retrieving per-semaphore information (Example)

Example showing how to **get information about a semaphore** of the semaphore set

```
key_t key = //... (generate a key in some way, i.e. with ftok)
// get, or create, the semaphore set
int semid = semget(key, 10, IPC_CREAT | S_IRUSR | S_IWUSR);
// ...
// get information about the first semaphore of the semaphore set
printf("Sem:%d getpid:%d getncnt:%d getzcnt:%d\n",
semid,
semctl(semid, 0, GETPID, NULL),
semctl(semid, 0, GETNCNT, NULL),
semctl(semid, 0, GETZCNT, NULL));
```

Once returned, the semaphore may already have changed state!



Semaphores

Other Operations



Semaphore Operations

The `semop` system call performs one or more operations (wait (P) and signal (V)) on semaphores.

```
#include <sys/sem.h>

// Returns 0 on success, or -1 on error
int semop(int semid, struct sembuf *sops, unsigned int nsops);
```

The `sops` argument is a pointer to an array that contains a sorted sequence of operations to be performed atomically, and `nsops` (> 0) gives the size of this array. The elements of the `sops` array are structures of the following form:

```
struct sembuf {
    unsigned short sem_num; /* Semaphore number */
    short sem_op;           /* Operation to be performed */
    short sem_flg;         /* Operation flags */
};
```



Semaphore Operations

The `sem_num` field identifies the semaphore within the set upon which the operation is to be performed. The `sem_op` field specifies the operation to be performed:

- `sem_op > 0`: value of `sem_op` is added to the value of the `sem_num`-th semaphore.
- `sem_op = 0`: the value of the `sem_num`-th semaphore is checked to see whether it currently equals 0. If it doesn't, the calling process is blocked until the semaphore is 0.
- `sem_op < 0`: decrease the value of the `sem_num`-th semaphore by the amount specified in `sem_op`. it blocks the calling process until the semaphore value has been increased to a level that permits the operation to be performed without resulting in a negative value.



Semaphore Operations

When a `semop(...)` call blocks, the process remains blocked until one of the following occurs:

- Another process modifies the value of the semaphore such that the requested operation can proceed.
- A signal interrupts the `semop(...)` call. In this case, the error `EINTR` results.
- Another process deletes the semaphore referred to by `semid`. In this case, `semop(...)` fails with the error `EIDRM`.

We can prevent `semop(...)` from blocking when performing an operation on a particular semaphore by specifying the `IPC_NOWAIT` flag in the corresponding `sem_flg` field. In this case, if `semop(...)` would have blocked, it instead fails with the error `EAGAIN`.



Semaphore Operations

Example showing how to initialize an array of sembuf operations

```
struct sembuf sops[3];

sops[0].sem_num = 0;
sops[0].sem_op = -1; // subtract 1 from semaphore 0
sops[0].sem_flg = 0;

sops[1].sem_num = 1;
sops[1].sem_op = 2; // add 2 to semaphore 1
sops[1].sem_flg = 0;

sops[2].sem_num = 2;
sops[2].sem_op = 0; // wait for semaphore 2 to equal 0
sops[2].sem_flg = IPC_NOWAIT; // but don't block if operation cannot be
    performed immediately
```



Semaphore Operations

Example showing how to perform operations on a semaphore set

```
struct sembuf sops[3];

// .. see the previous slide to initialize sembuf

if (semop(semid, sops, 3) == -1) {
    if (errno == EAGAIN) // Semaphore 2 would have blocked
        printf("Operation would have blocked\n");
    else
        errExit("semop"); // Some other error
}
```

